

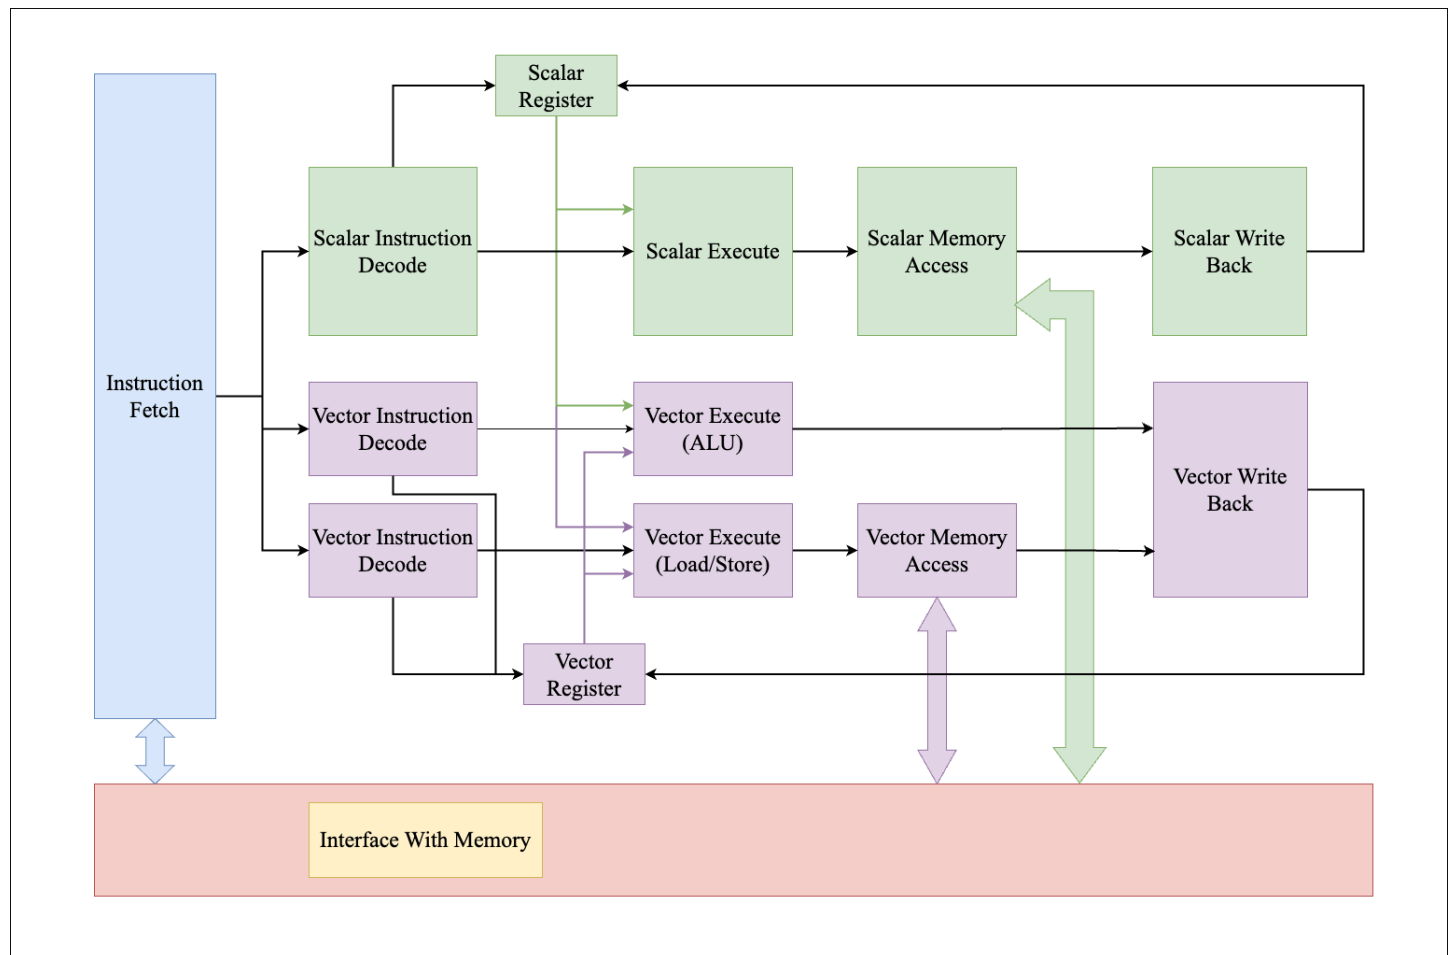
Lab-4 Design of a RISC-V Processor

Introduction

The lab aims to design a RISC-V processor supporting MAC (Multiply Accumulate) operation.

The lab introduces a multi-issue processor with a scalar datapath and two vector datapaths.

By completing the Verilog codes and designing the corresponding assembly codes, you will have a deep understanding of the computer architecture.



Background

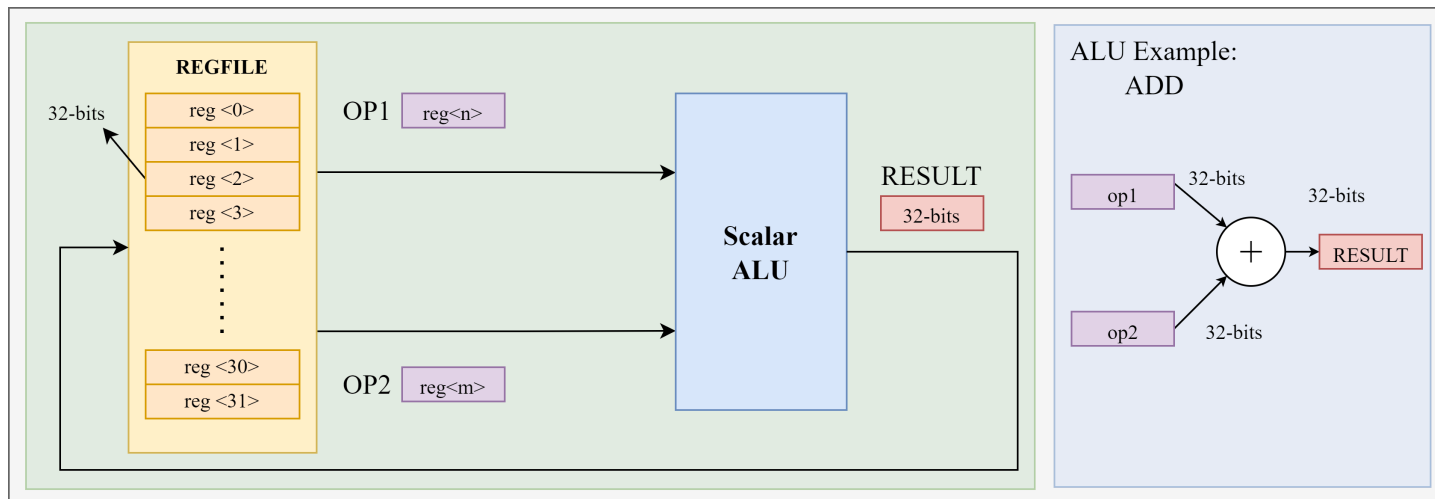
Vector Processor

We can use **SIMD** to achieve data parallelism. With the rise of applications such as multimedia, big data, and artificial intelligence, it has become increasingly important to endow processors with SIMD processing power. As these applications have a large number of fine-grained, homogeneous, and independent data operations that SIMD is inherently suited to handle.

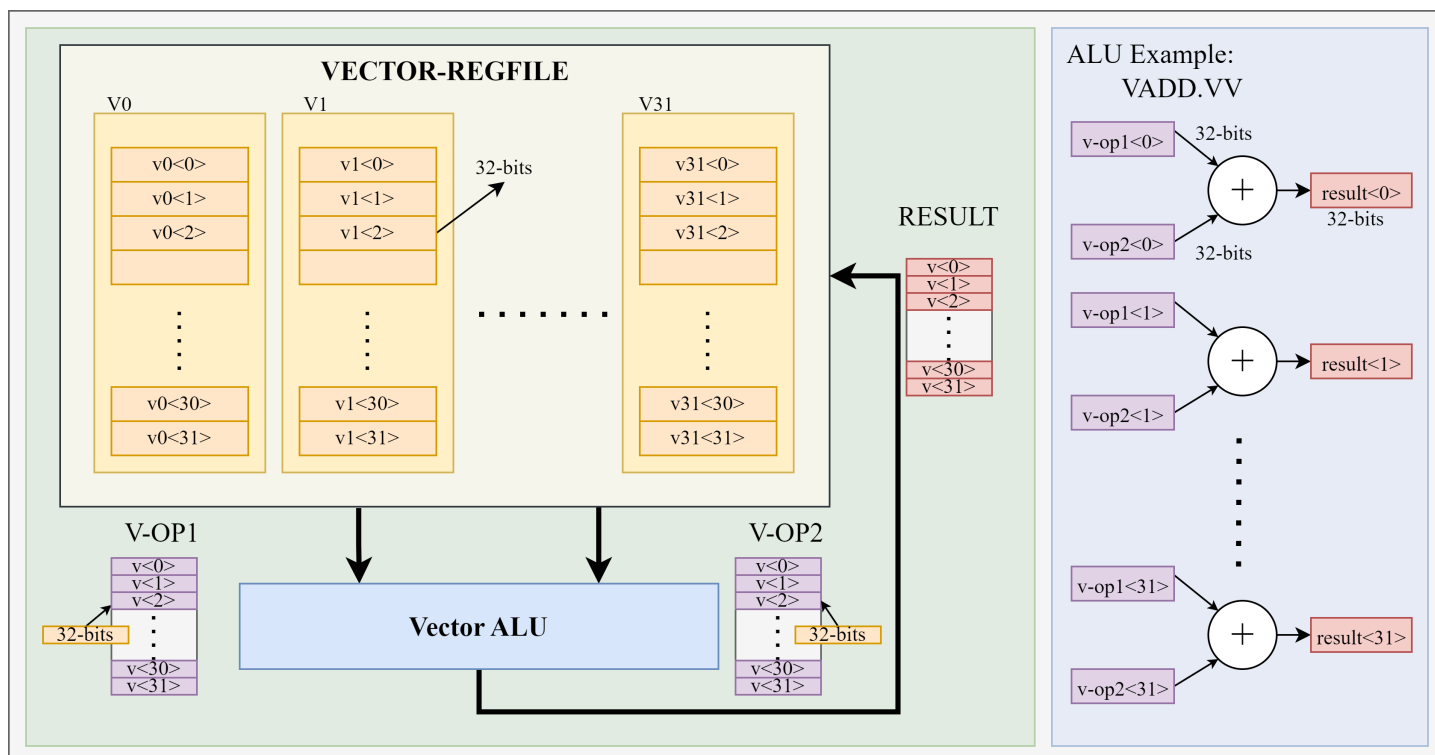
The structure of a vector processor fits well with the problem of parallel computation of large amounts of data. A vector processor has multiple ALUs that are capable of performing the same operation many times at the same time. The basic idea of the vector architecture is to collect data elements from memory, put them into a large set of registers, then operate on them using a pipelined execution unit, and finally write the result back to memory.

The key feature of vector architecture is a set of vector registers.

The next picture shows a simplified scalar processor's datapath.



The next picture shows a simplified vector processor's datapath. The scalar operand in scalar architecture is expanded to vector operand in vector architecture. And the relative REGFILE and ALU are also vectorized.

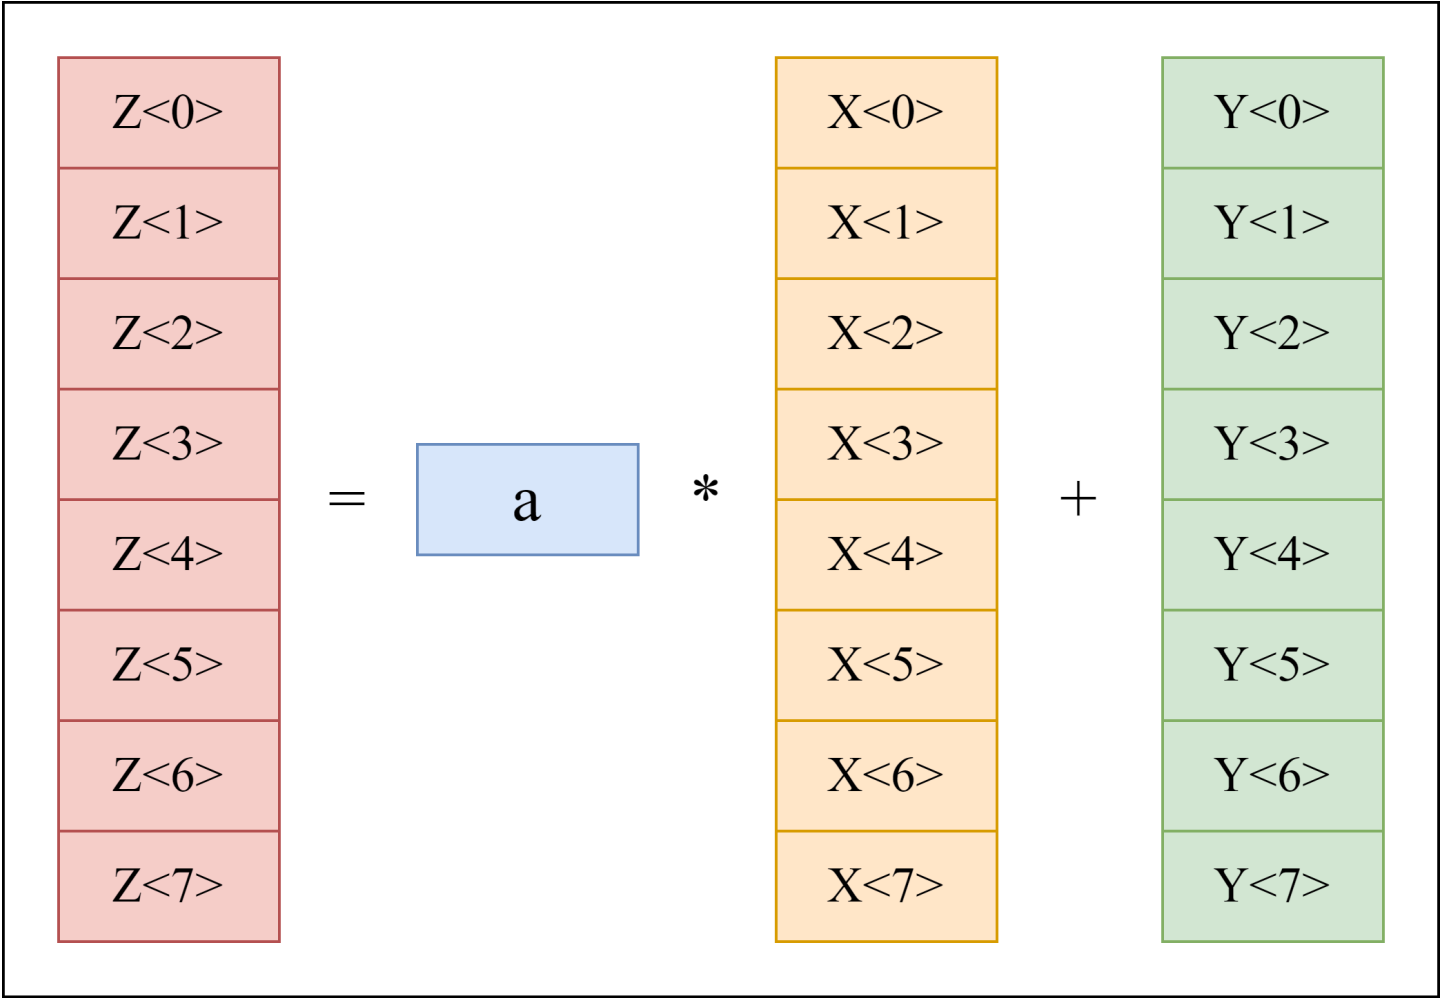


Vector VS Scalar

In order to visualize the characteristics of vector processors and conventional processors, we provide one case study. We want to use a vector processor and a conventional processor to perform the following operations, respectively.

$$Z = a \times X + Y$$

Here X , Y , Z are 8-dimensional vectors, and each element of the vector is one 32-bit integer data; a is a 32-bit integer scalar.



The address of a in memory is in register $x4$, the base address of X in memory is in register $x5$, the base address of Y in memory is in register $x6$, and the base address of Z in memory is in register $x7$.

variables	a	X	Y	Z
base address register	x4	x5	x6	x7

Scalar Processor

The assembly codes based on the RISC-V instruction set are shown below.

```
addi    x1,    $zero, 1      ; set x1 = 1
lw      x11,   0(x4)         ; load scalar a
addi    x12,   $zero, 8      ; upper bound of what to load
loop:
lw      x13,   0(x5)         ; load X[i]
mul     x13,   x13, x11       ; a x X[i]
lw      x14,   0(x6)         ; load Y[i]
add     x14,   x14, x13       ; a x X[i] + Y[i]
sw      x14,   0(x7)         ; store Z[i]
addi    x5,    x5,    4       ; increment index to x
addi    x6,    x6,    4       ; increment index to y
addi    x7,    x7,    4       ; increment index to z
sub     x12,   x12,   x1      ; x12 = x12 - 1
bne     x12,   $zero, loop    ; check if done
```

Vector Processor

The assembly codes based on the RISC-V Vector-Extension instruction set are shown below.

```
lw      x11,    0(x4)      ; load scalar a
vle32.v v13,    0(x5)      ; load vector X
vmul.vx v14,    v13,    x11 ; a x X
vle32.v v15,    0(x6)      ; load vector Y
vadd.vv v16,    v14,    v15 ; Z = a x X + Y
sle32.v v16,    0(x7)      ; store Z
```

Comparison

By comparing two assembly code implementations, we can discover that the number of instructions of the vector processor is much less than the number of instructions of the scalar version. This is mainly due to the vector processor's ability to compute multiple data in parallel, eliminating the need to use for loop.

Multi-Issue Processor

In computer architecture, multi-issue processors enable a processor to launch multiple instructions in a single clock cycle. There are mainly two categories of multi-issue processors: static multi-issue and dynamic multi-issue.

Dynamic multi-issue is implemented in hardware where the processor dynamically decides at runtime which instructions to issue, resolving dependencies and reordering instructions on-the-fly. This approach offers flexibility but adds hardware complexity and power consumption.

Static multi-issue, on the other hand, relies on the compiler to analyze and schedule instructions. It simplifies hardware but requires sophisticated compiler techniques and can be less adaptable to runtime conditions.

In this lab, we only focus on the static multi-issue processors.

The following figure gives an instruction processing timing diagram for a 2-issue pipelined processor. ALU/Branch instructions and Load/Store instructions can be seen to be processed in parallel at the same time. In order to support this operation, the hardware needs to have separate data paths for each of them.

Instruction type	Pipe stages							
ALU or branch instruction	IF	ID	EX	MEM	WB			
Load or store instruction	IF	ID	EX	MEM	WB			
ALU or branch instruction		IF	ID	EX	MEM	WB		
Load or store instruction		IF	ID	EX	MEM	WB		
ALU or branch instruction			IF	ID	EX	MEM	WB	
Load or store instruction			IF	ID	EX	MEM	WB	
ALU or branch instruction				IF	ID	EX	MEM	WB
Load or store instruction				IF	ID	EX	MEM	WB

To help understand the concepts, we rewrite the previous assembly code for a single-issue scalar processor as two-issue assembly code. Note that the code presented on the same line is executable in parallel.

```
addi    x1,    $zero, 1      ; lw      x11,    0(x4)      ; set x1 = 1 & load scalar a
addi    x12,   $zero, 8      ; nop                          ; upper bound of what to load
loop:
nop      ; lw      x13,    0(x5)      ; load X[i]
mul      x13,   x13,    x11      ; lw      x14,    0(x6)      ; a x X[i] & load Y[i]
add      x14,   x14,    x13      ; nop                          ; a x X[i] + Y[i]
addi     x5,    x5,    4          ; sw      x14,    0(x7)      ; increment index to x & store Z[i]
addi     x6,    x6,    4          ; nop                          ; increment index to y
addi     x7,    x7,    4          ; nop                          ; increment index to z
sub      x12,   x12,    x1        ; nop                          ; x12 = x12 - 1
bne      x12,   $zero, loop      ; nop                          ; check if done
```

RISC-V

RISC-V is an open-source instruction set architecture (ISA) based on the principles of Reduced Instruction Set Computing (RISC), with V denoting the fifth generation of RISC.

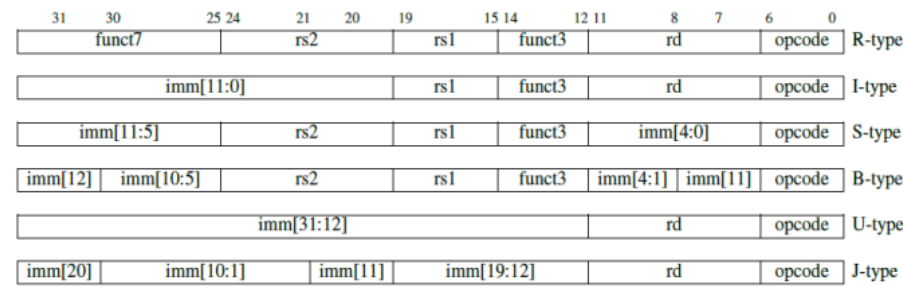
RISC-V instruction set includes RV32I and RV32M, RV32F, RV32D, RV32A, RV32V extensions. The instruction set RV32I is a fixed basic integer instruction set. RV32V is the vector extension.

In this lab, RV32I, RV32M and RV32V instruction sets are used.

RV32I

The following figure shows six basic instruction formats of RV32I:

- 1. R-type instructions for register-register operations
- 2. I-type instructions for short immediate and load operations
- 3. S-type instructions for store operations
- 4. B-type instructions for conditional jump operations
- 5. U-type instructions for long immediate
- 6. J-type instructions for unconditional jumps



The following figure shows the instructions of each type in RV32I.

31	25 24	20 19	15 14	12 11	7 6	0	
imm[31:12]				rd	0110111		U lui
imm[31:12]				rd	0010111		U auipc
imm[20 10:1 11 19:12]				rd	1101111		J jal
imm[11:0]		rs1	000	rd	1100111		I jalr
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011		B beq
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011		B bne
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011		B blt
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011		B bge
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011		B bltu
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011		B bgeu
imm[11:0]		rs1	000	rd	0000011		I lb
imm[11:0]		rs1	001	rd	0000011		I lh
imm[11:0]		rs1	010	rd	0000011		I lw
imm[11:0]		rs1	100	rd	0000011		I lbu
imm[11:0]		rs1	101	rd	0000011		I lhu
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011		S sb
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011		S sh
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011		S sw
imm[11:0]		rs1	000	rd	0010011		I addi
imm[11:0]		rs1	010	rd	0010011		I slti
imm[11:0]		rs1	011	rd	0010011		I sltiu
imm[11:0]		rs1	100	rd	0010011		I xori
imm[11:0]		rs1	110	rd	0010011		I ori
imm[11:0]		rs1	111	rd	0010011		I andi
0000000	shamt	rs1	001	rd	0010011		I slli
0000000	shamt	rs1	101	rd	0010011		I srli
0100000	shamt	rs1	101	rd	0010011		I srai
0000000	rs2	rs1	000	rd	0110011		R add
0100000	rs2	rs1	000	rd	0110011		R sub
0000000	rs2	rs1	001	rd	0110011		R sll
0000000	rs2	rs1	010	rd	0110011		R slt
0000000	rs2	rs1	011	rd	0110011		R sltu
0000000	rs2	rs1	100	rd	0110011		R xor
0000000	rs2	rs1	101	rd	0110011		R srl
0100000	rs2	rs1	101	rd	0110011		R sra
0000000	rs2	rs1	110	rd	0110011		R or
0000000	rs2	rs1	111	rd	0110011		R and
0000	pred	succ	00000	000	00000	0001111	I fence
0000	0000	0000	00000	001	00000	0001111	I fence.i
000000000000			00000	00	00000	1110011	I ecall
000000000000			00000	000	00000	1110011	I ebreak
csr		rs1	001	rd	1110011		I csrrw
csr		rs1	010	rd	1110011		I csrrs
csr		rs1	011	rd	1110011		I csrrc
csr		zimm	101	rd	1110011		I csrrwi
csr		zimm	110	rd	1110011		I csrrsi
csr		zimm	111	rd	1110011		I csrrci

RV32M

RV32M adds integer multiplication and division instructions to RV32I, including instructions for signed and unsigned integers, divide (div) and divide unsigned (divu), which put the quotient into the target register. In a few cases, the programmer needs the remainder rather than the quotient, so RV32M provides remainder (rem) and remainder unsigned (remu), which writes the remainder to the target register instead of the quotient.

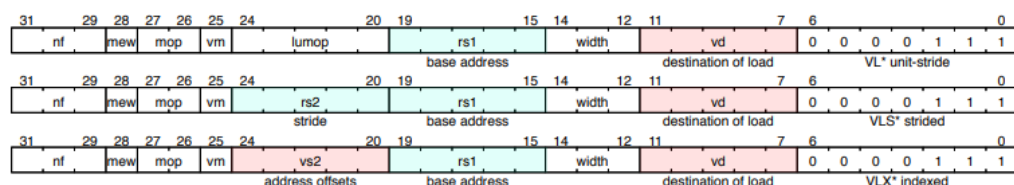
To correctly obtain a signed or unsigned 64-bit product, RISC-V comes with four multiplication instructions. To get the integer 32-bit product (the lower 32 bits of 64 bits) use the mul instruction. To get the high 32 bits, use the mulh instruction for signed operands, the mulhu instruction for unsigned operands and the mulhsu instruction for signed multiplier and unsigned multiplicand.

31	25	24	20	19	15	14	12	11	7	6	0	
0000001		rs2		rs1		000		rd		0110011		R mul
0000001		rs2		rs1		001		rd		0110011		R mulh
0000001		rs2		rs1		010		rd		0110011		R mulhsu
0000001		rs2		rs1		011		rd		0110011		R mulhu
0000001		rs2		rs1		100		rd		0110011		R div
0000001		rs2		rs1		101		rd		0110011		R divu
0000001		rs2		rs1		110		rd		0110011		R rem
0000001		rs2		rs1		111		rd		0110011		R remu

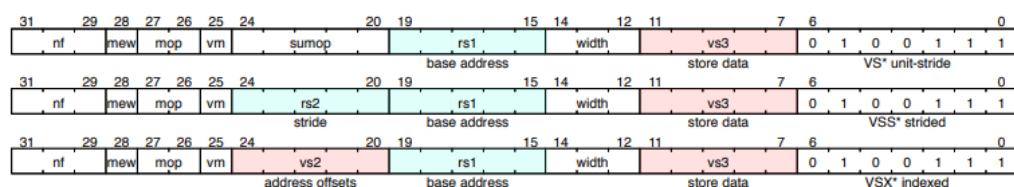
RV32V

The opcodes of the instructions in the vector extension either are the same as LOAD-FP and STORE-FP or use OP-V.

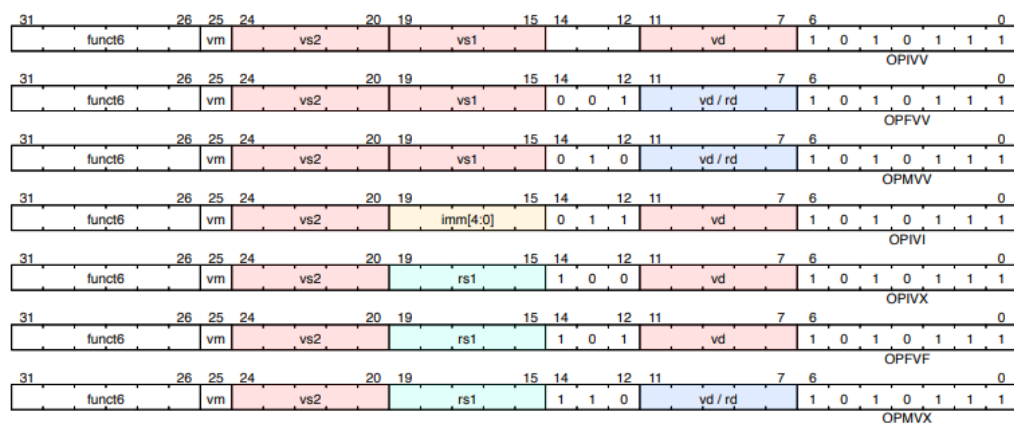
Format for Vector Load Instructions under LOAD-FP major opcode



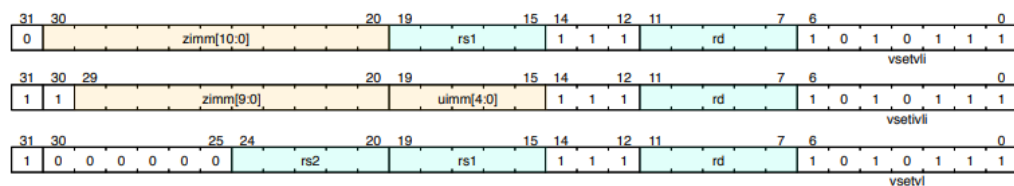
Format for Vector Store Instructions under STORE-FP major opcode



Formats for Vector Arithmetic Instructions under OP-V major opcode



Formats for Vector Configuration Instructions under OP-V major opcode



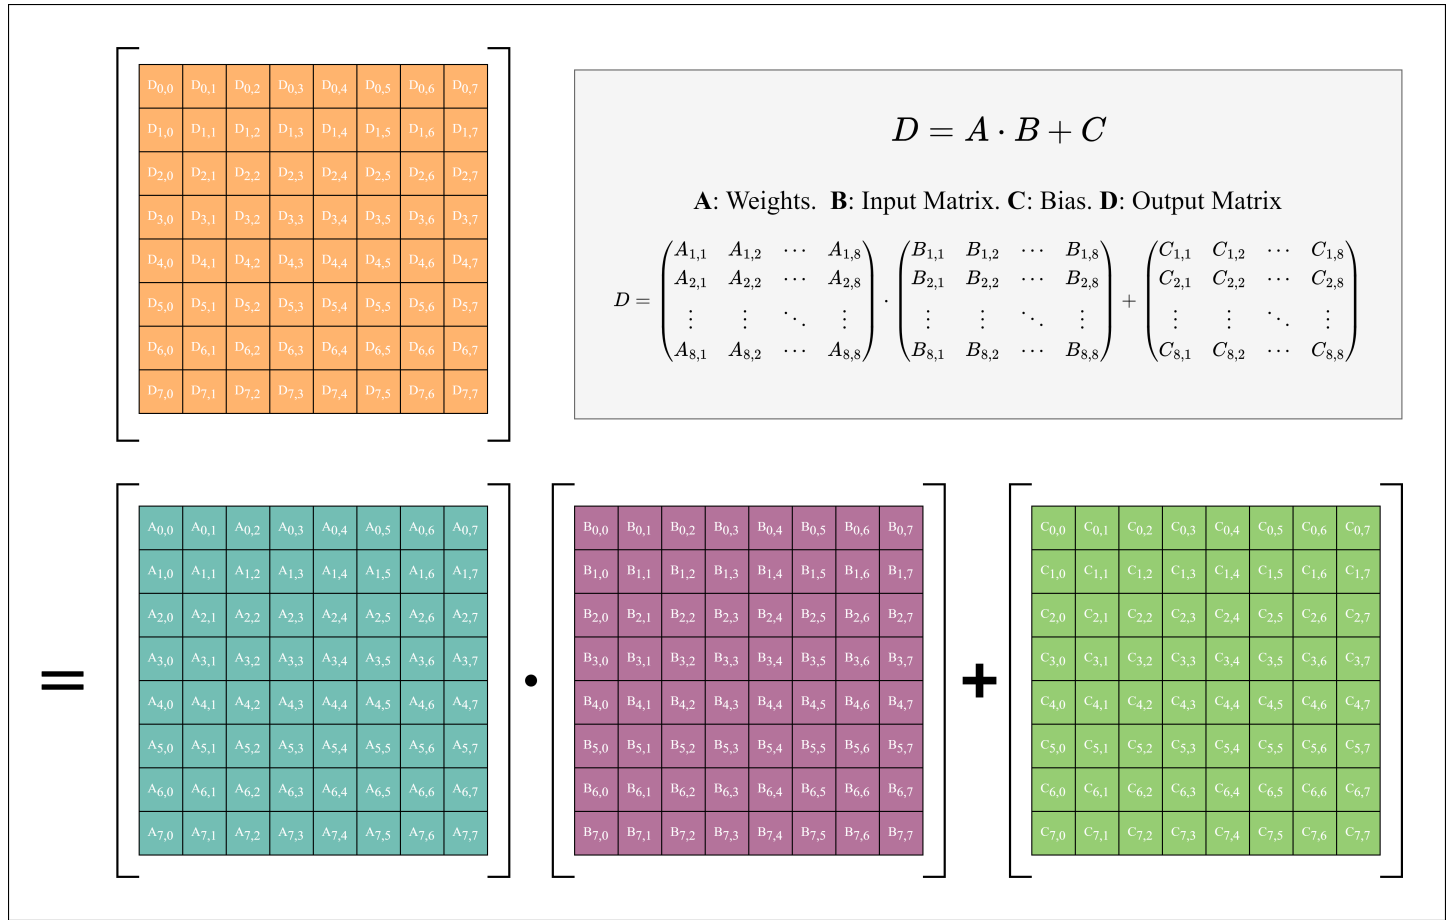
Category	operation	Operands	Type of scalar operand
OPI-VV	integer operation	vector-vector	N/A
OPF-VV	floating-point operation	vector-vector	N/A
OPM-VV	mask operation	vector-vector	N/A
OPI-VI	integer operation	vector-immediate	imm[4:0]
OPI-VX	integer operation	vector-scalar	GPR(general purpose registers) x register rs1
OPI-VF	integer operation	vector-scalar	FP f register rs1
OPM-VX	mask operation	vector-scalar	GPR x register rs1

In this lab, we only focus on the OPIVV, OPIVI and OPIVX.

Lab Goals

The purpose of the lab is to design a multi-issue RISC-V processor, which can perform the MAC operation. The expected operation is shown below. The dimensions of all matrices are 8 by 8 and each element in the matrices is a 32-bit data. Matrix-A is the weight matrix. Matrix-B is the input matrix. Matrix-C is the bias matrix. Matrix-D is the output matrix.

For simplicity, we assume no data overflow will occur.

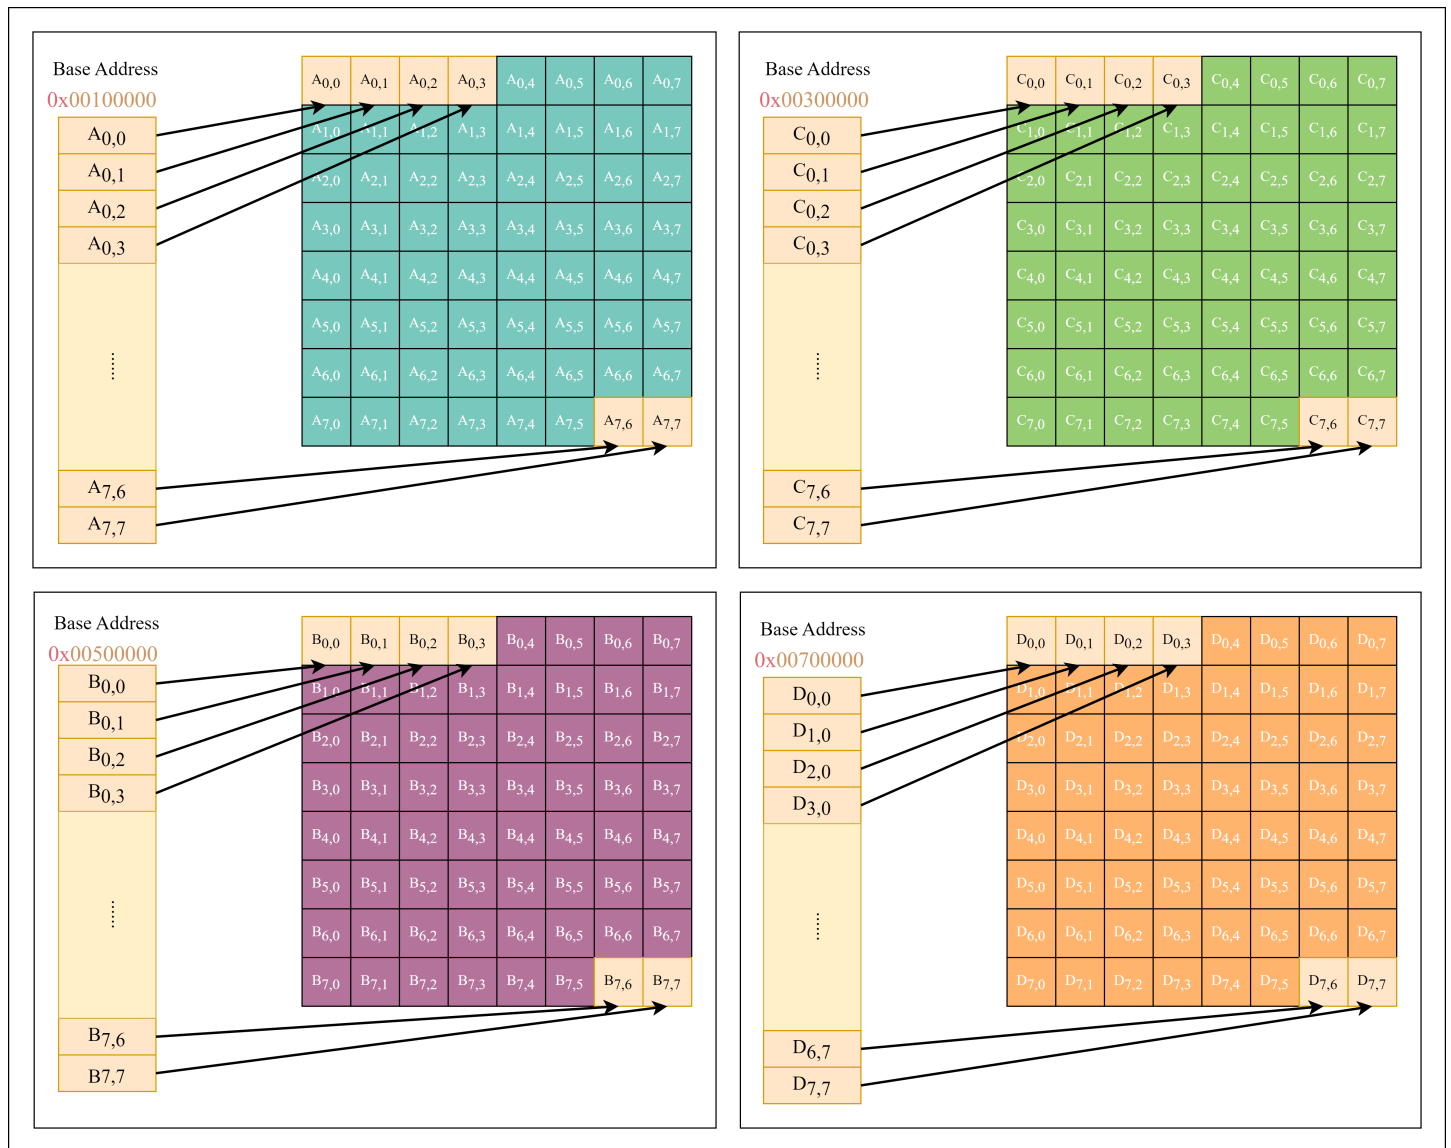


Lab Setup

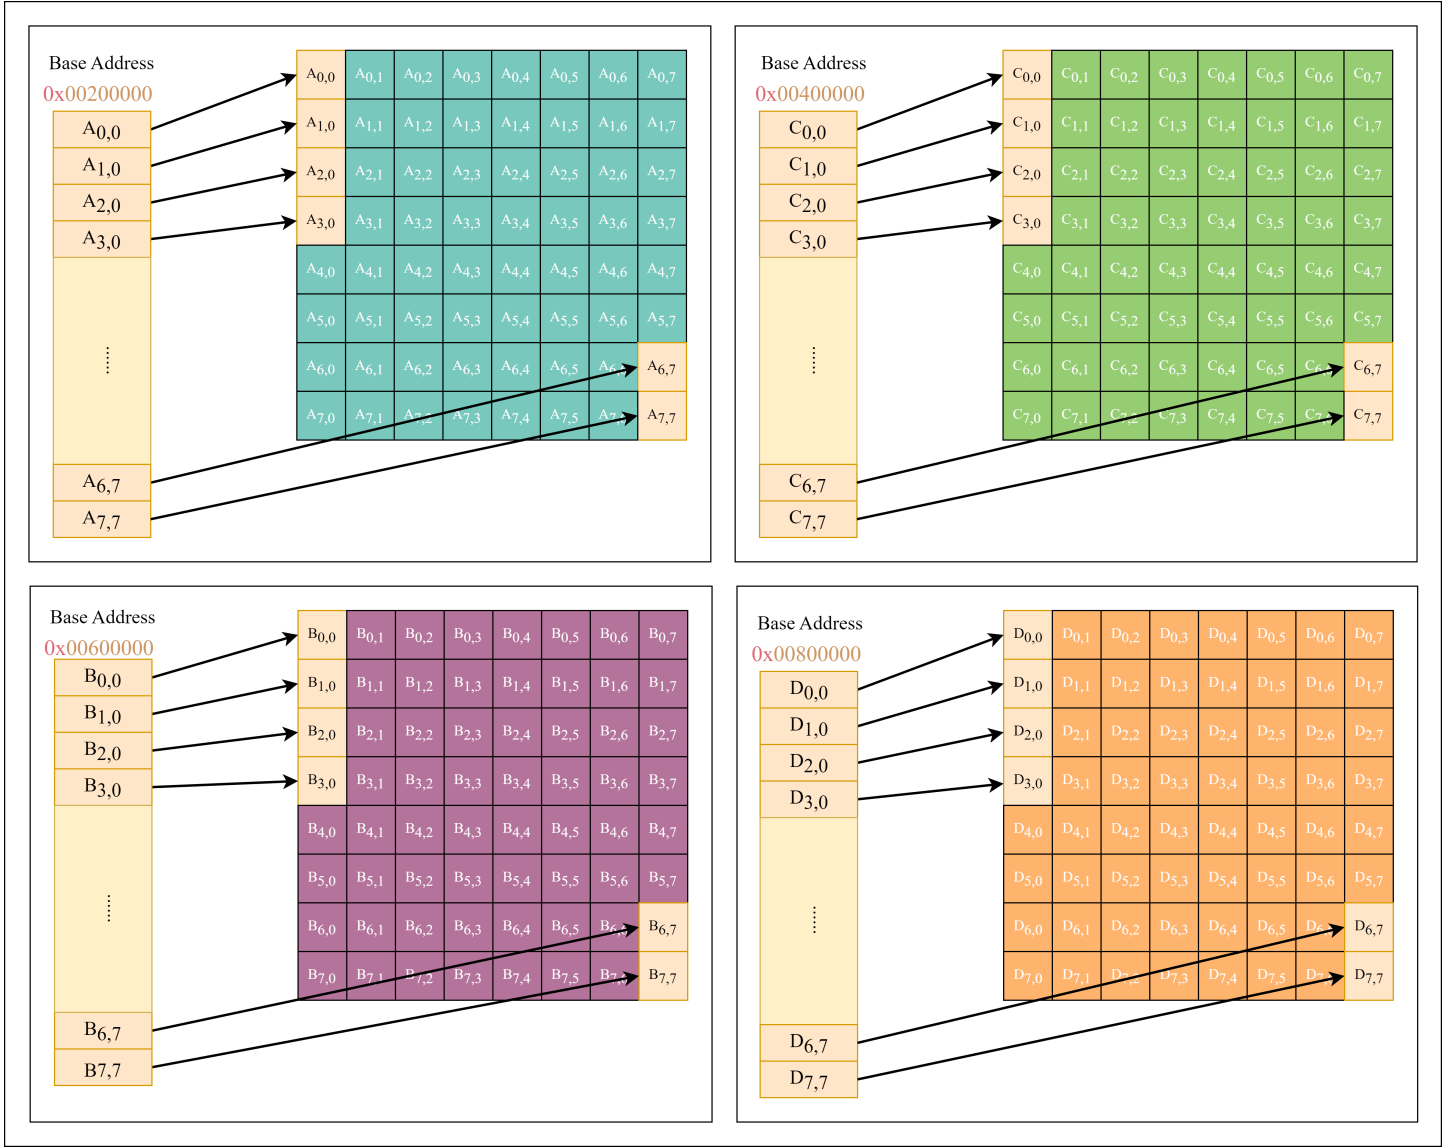
Data Mapping

The mapping of the data of the matrix to the addresses of the memory are shown in the following figures. For all the matrixes involved, we provide two ways of data mapping.

One is to store the matrices in row-major.



The other is to store the transposed matrices in column-major.



The base address is the first element's address of the data block. The detailed base address information is shown below.

Data	Base Address
Instructions	0x80000000
Matrix-A	0x80100000
Matrix-A.T	0x80200000
Matrix-B	0x80300000
Matrix-B.T	0x80400000
Matrix-C	0x80500000
Matrix-C.T	0x80600000
Matrix-D	0x80700000
Matrix-D.T	0x80800000

Memory Interface

To reduce the challenge of accessing memory, we use 3 simplified memory access models.

Caution : The memory interface has been greatly simplified here to reduce the difficulty of the experiment. The memory can read out data in one clock cycle (processor clock) during simulation. But in practice, the memory's read and write speed are much lower than the processor's main frequency. In addition, the memory and the processor are often connected through the bus, and instruction storage and data storage do not necessarily use two sets of read-out interfaces. Thus, the processor needs to arbitrate whether to read data or the instructions.

In this lab, the INST_RAMHelper is used to get the instructions from the memory.

```
module INST_RAMHelper(  
    input          clk,  
    input          ren,    // read enable  
    input  [31:0]  raddr,  // read address  
    output [31:0]  rdata   // read data  
);
```

The SCALAR_RAMHelper is used to get the scalar data from the memory.

```
module SCALAR_RAMHelper(  
    input          clk,  
    input          ren,    // read enable  
    input  [31:0]  raddr,  // read address  
    output [31:0]  rdata,  // read data  
    input          wen,    // write enable  
    input  [31:0]  waddr,  // write address  
    input  [31:0]  wdata,  // write data  
    input  [31:0]  wmask   // write mask  
);
```

The Vector_RAMHelper is used to get the vector data from the memory.

```
module VECTOR_RAMHelper(  
    input          clk,  
    input          ren,    // read enable  
    input  [31 :0]  raddr,  // read address  
    output [255:0]  rdata,  // read data  
    input          wen,    // write enable  
    input  [31 :0]  waddr,  // write address  
    input  [255:0]  wdata,  // write data  
    input  [255:0]  wmask   // write mask  
);
```

Data Access

The minimum stride of access memory is 1 Byte (8 bits).

For 32-bit data access, the access address's step is 4. For example, data A and data B are both 32-bit data and stored in adjacent locations. Data A's address is 0x0000_1000 and data B's address is 0x0000_1004.

RISC-V Vector Extension

Vector Data Width Setting

The vector extended RISC-V processor contains 32 vector registers, `x0~x31` .

`VLEN` : represents a fixed bit width for each vector register.

`SEW` : represents the bit width of the selected vector element. It is controlled by register `vsew[2:0]`.

Taking `VLEN=128bits` as an example, the correspondence between the number of elements contained in each vector register and `SEW` is shown in the following table.

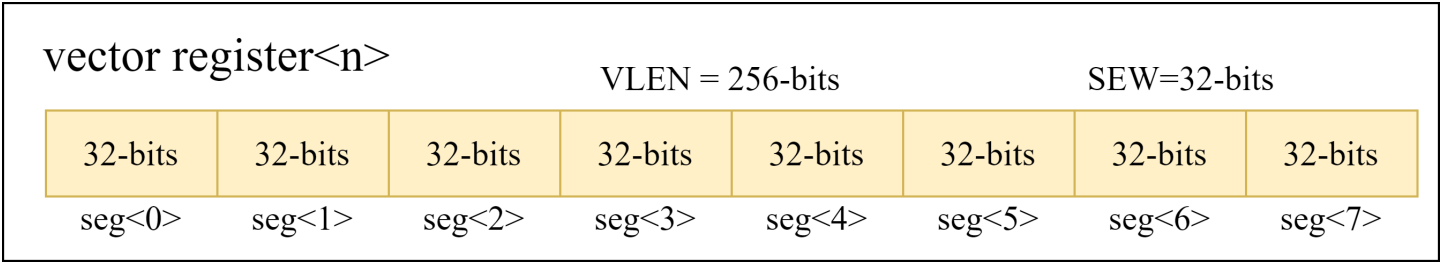
VELN	SEW	Elements Per Vector Register
128	64	2
128	32	4
128	16	8
128	8	16

`LMUL` : represents the vector length multiplier. If greater than 1, it represents the default number of vector registers.

`VLMAX` : represents the maximum number of vector elements that a vector instruction can operate on. `VLMAX=LMUL*VLEN/SEW` .

!!! IMPORTANT !!!

To simplify the experiment, we set all the above parameters to fixed values.



Mask Setting

In the arithmetic and access instructions, you can choose whether to use the mask or not, and whether to enable the function is controlled by the `vm` bit in the instruction.

!!! IMPORTANT !!!

In this experiment, the mask is not used by default. And `vm` is set to 1 in all instructions to disable the mask function.

Load/Store Setting

There are 3 different access modes for RISC-V vector extensions: `unit-stride`, `strided`, `indexed`.

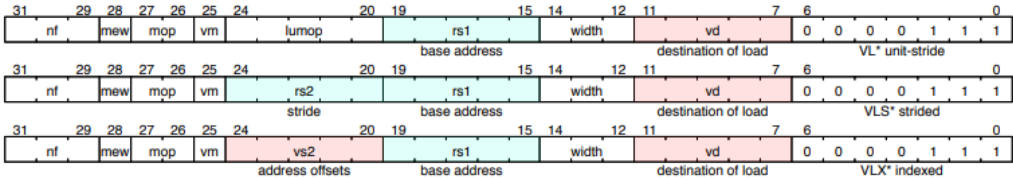
The `unit-stride` operation accesses consecutive elements stored in memory starting at the base effective address.

The `strided` operation accesses the first memory element at the base effective address, and then accesses subsequent elements at the address increment given by the byte offset contained in the `x` register specified by `rs2`.

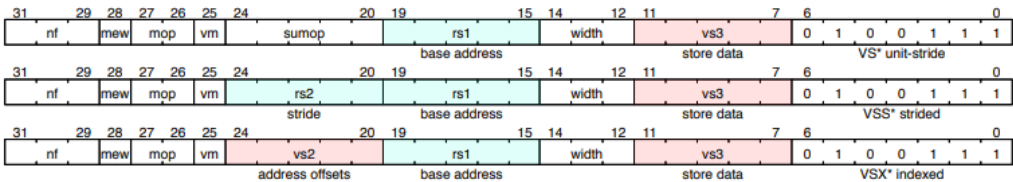
The `indexed` operation adds the vector offset specified by `rs2` to the base effective address to obtain the effective address of each element.

The following images describe the format of the access command and its specific meaning.

Format for Vector Load Instructions under LOAD-FP major opcode



Format for Vector Store Instructions under STORE-FP major opcode



Field	Description
rs1[4:0]	specifies x register holding base address
rs2[4:0]	specifies x register holding stride
vs2[4:0]	specifies v register holding address offsets
vs3[4:0]	specifies v register holding store data
vd[4:0]	specifies v register destination of load
vm	specifies whether vector masking is enabled (0 = mask enabled, 1 = mask disabled)
width[2:0]	specifies size of memory elements, and distinguishes from FP scalar
mew	extended memory element width. See Vector Load/Store Width Encoding
mop[1:0]	specifies memory addressing mode
nf[2:0]	specifies the number of fields in each segment, for segment load/stores
lumop[4:0]/sumop[4:0]	are additional fields encoding variants of unit-stride instructions

!!! IMPORTANT !!!

In order to simplify the experiment, this experiment **only needs to support the access mode of `unit-stride`**. In addition, the `nf`, `mew`, `mop`, and `lumop` bits of the access instruction can be set to the default value of 0.

Assembly Code

The assembly code needs to be translated into machine code by the translator in order to be recognized and executed by the processor. A simple translator is provided in the experimental environment to translate the assembly code into a bin file. Using this translator requires the assembly code to be written in the specified format.

If you are not familiar with the assembly code, you can try to use the RISC-V assembler and runtime simulator.

[TheThirdOne/rars: RARS -- RISC-V Assembler and Runtime Simulator \(github.com\)](https://github.com/TheThirdOne/rars)

The file `demo.asm` in the `asm` folder shows an example code that calls all the used instructions in the lab. Please refer to the format of this code to write assembly code.

The scalar registers' names are `x0-x31`. And `x0` can be replaced with `zero`.

The vector registers' names are `vx0-vx31`. And `vx0` can be replaced with `vzero`.

```
; =====
; scalar instructions
; =====

; load matrix_A base-address 0x80100000 to the register x5
; ( the immediate number is already shifted 12-bits here )
lui    x5,    2148532224    ; x5 = 0x80100000

addi   x6,    x5,    4        ; x6 = x5+4
lw     x7,    0(x5)          ; x7 = A[0][0]
lw     x8,    4(x6)          ; x18 = A[0][2]

slti   x9,    x7,    64        ; x9 = (x7<64)
addi   x9,    zero,    1      ; x9 = 1
add    x10,   x9,    x9        ; x10 = x9+x9
and    x11,   x10,   x9        ; x11 = x10 & x9
mul    x12,   x10,   x10       ; x12 = x10 * x10
sll    x13,   x9,    1         ; x13 = x9<<1

addi   x9,    zero,    4        ; assign 4 to x9
addi   x10,   zero,    0        ; assign 0 to x10
loop:
addi   x10,   x10,    1        ; assign x10+1 to x10
mul    x12,   x12,   x10       ; assign x12*x10 to x12
blt    x10,   x9,    loop      ; if( x10<x9 ) jump to loop

sw     x12,   0(x5)          ; A[0][0] = x12
jal    x1,    label          ; jump to label

; =====
; vector instructions
; =====
label:
vle32.v vx2,   x5,    1        ; vx2 = A[0][0:7]

vadd.vi vx3,   vzero, 1,    1  ; vx3[i] = 1
vadd.vx vx3,   vzero, x12, 1   ; vx3[i] = x12
vmul.vv vx4,   vx2,   vx3, 1   ; vx4[i] = vx2[i] * vx3[i]

vse32.v vx4,   x5,    1        ; A[0][0:7] = vx4
```

Detailed Tasks

Task1.1 - Single-Issue Processor

You need to complete the instructions in the table below and pass the test. Here `sext` means `signed extended` and `Mem` means `Memory`. The logic operation (`&` and `)` is bit-wise.

Type	Instruction	Format	Implementation
R	mul	mul rd, rs1, rs2	$x[rd] = x[rs1] * x[rs2]$
R	add	add rd, rs1, rs2	$x[rd] = x[rs1] + x[rs2]$
R	and	and rd, rs1, rs2	$x[rd] = x[rs1] \& x[rs2]$
R	sll	sll rd, rs1, rs2	$x[rd] = x[rs1] \ll x[rs2]$
I	addi	addi rd, rs1, imm	$x[rd] = x[rs1] + sext(immediate)$
I	slti	slti rd, rs1, imm	$x[rd] = x[rs1] < sext(immediate)$
I	lw	lw rd, offset(rs1)	$x[rd] = Mem[x[rs1] + sext(offset)]$
S	sw	sw rs2, offset(rs1)	$Mem[x[rs1] + sext(offset)] = x[rs2]$
B	blt	blt rs1, rs2, offset	if $(x[rs1] < x[rs2])$ $pc += sext(offset)$
U	lui	lui rd,imm	$x[rd] = immediate[31:12] \ll 12$
J	jal	jal rd,offset	$x[rd] = pc+4; pc += sext(offset)$

- To better understand the meaning of these instructions, you can refer to [this](#).
- To get the corresponding instruction format, you can refer to [this](#).
- The code templates have been provided for the implementation of several instructions.
- HINT: Finish the modules under /src/vsrc/components/single_issue. Top module (/src/vsrc/rvcpu/rvcpu_single_issue.v) is already given.

Task1.2 - Assembly Code 1

After completing Task 1.1, we have been able to use these instructions to perform some useful computation. Now, use these instructions to complete the [MAC](#) Operation.

HINT: Finish the code in /src/asm/task1_2.asm.

Task2.1 - Two-Issue Processor

Extend the scalar processor to a two-issue processor, with one datapath supporting scalar and another supporting vector operations. For vector operations, you need to complete the instructions in the table below and pass the test.

Type	Instruction	Format	Implementation
LOAD	vle32.v	vle32.v vd, rs1, vm	$vd = Mem[x[rs1]]$
STORE	vse32.v	vse32.v vs3, rs1, vm	$Mem[x[rs1]] = vs3$
OPIVV	vadd.vv	vadd.vv vd, vs2, vs1, vm	$vd[i] = vs2[i] + vs1[i]$
OPIVI	vadd.vi	vadd.vi vd, vs2, imm, vm	$vd[i] = vs2[i] + imm$
OPIVX	vadd.vx	vadd.vx vd, vs2, rs1, vm	$vd[i] = vs2[i] + x[rs1]$
OPIVV	vmul.vv	vmul.vv vd, vs2, vs1, vm	$vd[i] = vs2[i] * vs1[i]$
OPIVI	vmul.vi	vmul.vi vd, vs2, imm, vm	$vd[i] = vs2[i] * imm$
OPIVX	vmul.vx	vmul.vx vd, vs2, rs1, vm	$vd[i] = vs2[i] * x[rs1]$

The formats of relative instructions are shown below.

Instruction	Format
vle32.v	<div><div>31292827262524201915141211760</div><div>000000vm000000rs1110vd0000111</div><div>nf mew mop lumop width opcode (VL* unit-stride)</div></div>
vse32.v	<div><div>31292827262524201915141211760</div><div>000000vm000000rs1110vs30100111</div><div>nf mew mop sumop width opcode (VS* unit-stride)</div></div>

Instruction	Format
vadd.vv	<div><div><div>31262524201915141211760</div><div>000000vmvs2vs1000vd1010111</div><div>funct6funct3opcode (OPIVV)</div></div></div>
vmul.vv	<div><div><div>31262524201915141211760</div><div>100101vmvs2vs1000vd1010111</div><div>funct6funct3opcode (OPIVV)</div></div></div>
other	<div>vadd.vi , vadd.vx , vmul.vi , and vmul.vx only need to modify the funct3 . For OPIVI , the funct3 is 011 . For OPIVX , the funct3 is 100 .</div>

HINT: Finish the modules under /src/vsrc/components/two_issue and /src/vsrc/rvcpu/rvcpu_two_issue.v.

Task2.2 - Assembly Code 2

Vector processors can handle MAC operations more efficiently. Now, use the vector instructions to do this operation.

Note that you are only allowed to use instructions from the instruction set RV-32I and RV-32V. In other words, you are not allowed to use the mul instruction from RV-32M.

HINT: Finish the code in /src/asm/task2_2.asm.

Task3.1 - Three-Issue Processor

Split the datapath of the processor of Task2.1 into a three-issue processor with one scalar datapath, one vector alu datapath and one for vle and vse. Please refer to the figure of the top arhictecture at the beginning.

HINT: Finish the modules under /src/vsrc/components/vector_issue and /src/vsrc/rvcpu/rvcpu_three_issue.v.

Task3.2 - Assembly Code 2

Rewrite the assembly code to adapt it to a three-issue processor.

HINT: Finish the code in /src/asm/task3_2.asm.

Simulation Environment

File Tree

```
|-- build          --Folder containing the simulation files
|-- src
|   |-- asm        --Folder containing the assembly codes
|   |-- csrc       --Folder containing the cpp files
|   |-- vsrc       --Folder containing the verilog codes
|-- tools          --Folder containing the python scripts
|-- Makefile
```

Make Command

Run the simulation and test. IMG is your assembly code name. ISSUE_NUM determines the issue number (1 for Task1.x, 2 for Task2.x, 3 for Task3.x). You can modify the parameter IMG and ISSUE_NUM in Makefile or in your commands.

```
make run IMG=task1_2 ISSUE_NUM=1
make run IMG=task2_2 ISSUE_NUM=2
```

HINT: You can check your answer of Task x.1 by setting IMG = Taskx_1 and Task x.2 by setting IMG = Taskx_2.

There are 2 ways to clean the built files.

```
make clean # reserve the vcd file
make clear # clean all files in build_test
```

Grading

The total score (100%) is the sum of code (100%). Besides, you need to finish the content in `answer.sh` .

Code (100%)

- Task1 (30%)
 - Task1.1 (15%)
 - Task1.2 (15%)
- Task2 (30%)
 - Task2.1 (20%)
 - Task2.2 (10%)
- Task3 (40%)
 - Task3.1 (20%)
 - Task3.2 (20%)

Submission

Pack your lab4 folder and rename it with a new name `StudentNumber_StudentName_Lab4.zip` , and submit to Blackboard. The file structure should be like this.

12345678_张三_Lab4.zip

Reference Link

1. A reference video for last year's Lab4 is uploaded via bb.
2. [RISC-V 手册 一本开源指令集的指南](#)
3. [The RISC-V Instruction Set Manual Volume I: Unprivileged ISA](#)
4. [riscv-card](#)
5. [RV32I, RV64I Instructions](#)
6. [RISCV-V Online Doc](#)
7. [RISCV-V Document Pdf](#)